

## Some Important Required Libraries

```
In [57]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [58]: data=pd.read_csv("C:/Users/Dell/Downloads/Advertising.csv")
```

```
In [59]: data
```

```
Out[59]:
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
...	...	...	...	...	...
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

200 rows × 5 columns

```
In [60]: data.shape
```

```
Out[60]: (200, 5)
```

In [61]: data.info

Out[61]: <bound method DataFrame.info of Unnamed: 0 TV Radio Newspaper Sales  
 es  
 0 1 230.1 37.8 69.2 22.1  
 1 2 44.5 39.3 45.1 10.4  
 2 3 17.2 45.9 69.3 9.3  
 3 4 151.5 41.3 58.5 18.5  
 4 5 180.8 10.8 58.4 12.9  
 .. ...  
 195 196 38.2 3.7 13.8 7.6  
 196 197 94.2 4.9 8.1 9.7  
 197 198 177.0 9.3 6.4 12.8  
 198 199 283.6 42.0 66.2 25.5  
 199 200 232.1 8.6 8.7 13.4  
 [200 rows x 5 columns]>

In [62]: data.head(10)

Out[62]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
5	6	8.7	48.9	75.0	7.2
6	7	57.5	32.8	23.5	11.8
7	8	120.2	19.6	11.6	13.2
8	9	8.6	2.1	1.0	4.8
9	10	199.8	2.6	21.2	10.6

In [63]: data.describe()

Out[63]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000	200.000000
mean	100.500000	147.042500	23.264000	30.554000	14.022500
std	57.879185	85.854236	14.846809	21.778621	5.217457
min	1.000000	0.700000	0.000000	0.300000	1.600000
25%	50.750000	74.375000	9.975000	12.750000	10.375000
50%	100.500000	149.750000	22.900000	25.750000	12.900000
75%	150.250000	218.825000	36.525000	45.100000	17.400000
max	200.000000	296.400000	49.600000	114.000000	27.000000

```
In [64]: data.tail()
```

```
Out[64]:
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

```
In [65]: data.columns
```

```
Out[65]: Index(['Unnamed: 0', 'TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')
```

```
In [66]: data.dtypes
```

```
Out[66]: Unnamed: 0    int64
TV                float64
Radio             float64
Newspaper         float64
Sales             float64
dtype: object
```

```
In [67]: #Checking the null values in column
data.isnull().sum()
```

```
Out[67]: Unnamed: 0    0
TV                0
Radio             0
Newspaper         0
Sales             0
dtype: int64
```

```
In [68]: data.count()
```

```
Out[68]: Unnamed: 0    200
TV                200
Radio             200
Newspaper         200
Sales             200
dtype: int64
```

```
In [69]: corr_matrix = data.corr()
corr_matrix
```

```
Out[69]:
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
Unnamed: 0	1.000000	0.017715	-0.110680	-0.154944	-0.051616
TV	0.017715	1.000000	0.054809	0.056648	0.782224
Radio	-0.110680	0.054809	1.000000	0.354104	0.576223
Newspaper	-0.154944	0.056648	0.354104	1.000000	0.228299
Sales	-0.051616	0.782224	0.576223	0.228299	1.000000

```
In [70]: data['Sales'].value_counts()
```

```
Out[70]: 9.7      5
11.7      4
12.9      4
15.9      4
20.7      3
..
17.0      1
18.3      1
22.3      1
14.0      1
25.5      1
Name: Sales, Length: 121, dtype: int64
```

## Data Splitting

```
In [71]: # Split the data into input(x) and (y)variable
x = data[['TV', 'Radio', 'Newspaper']]
y = data['Sales']
```


```
In [72]: # Splitting the dataset into train and test sets
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [73]: print(x.shape, x_train.shape, x_test.shape)

(200, 3) (160, 3) (40, 3)
```


## Visualization Data

```
In [112]: sns.set(style='whitegrid')
plt.figure(figsize=(8,6))
sns.histplot(data['Sales'], bins=20, color='brown', kde=True)
plt.xlabel('Sales')
plt.ylabel('Frequency')
plt.show()
```



```
In [75]: plt.figure(figsize=(6,6))
sns.pairplot(data, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales', height=10,
plt.show())
```

<Figure size 600x600 with 0 Axes>

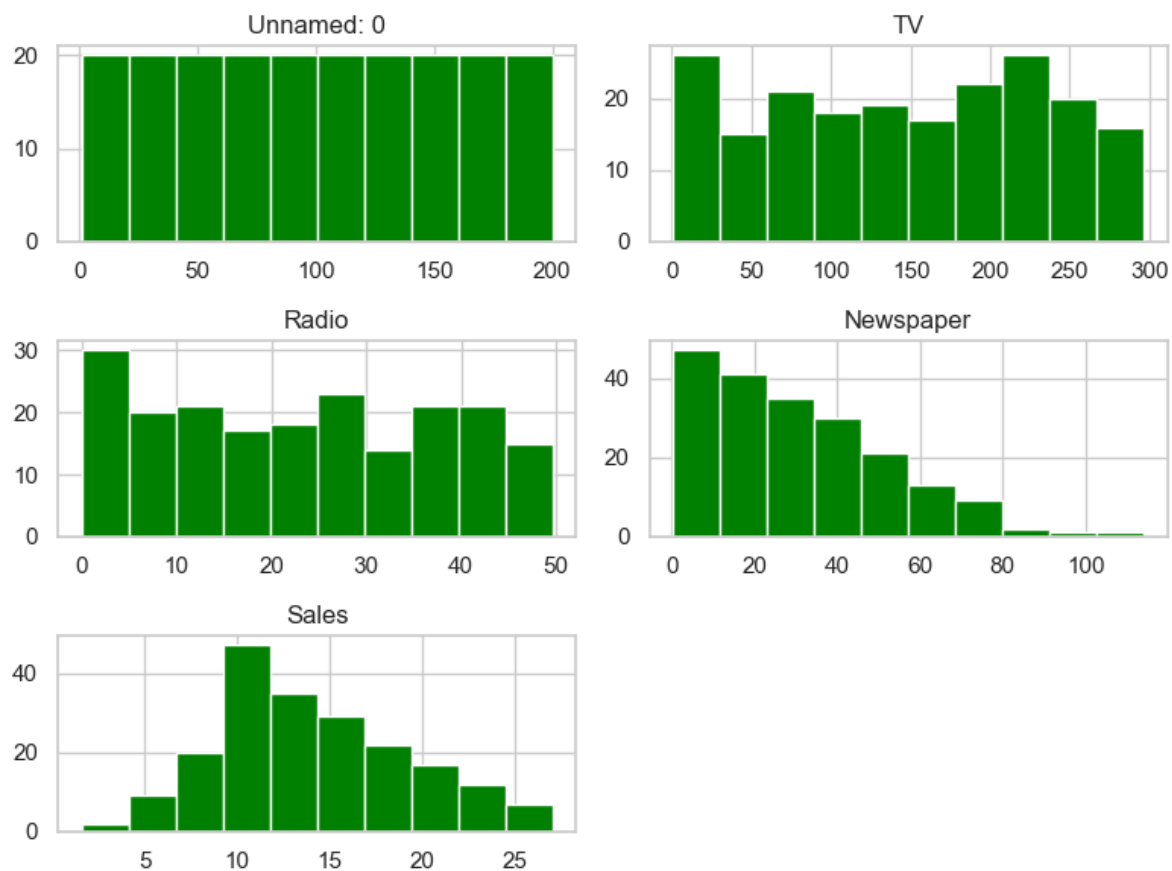


```
In [76]: sns.heatmap(df.isnull())
```


```
Out[76]: <AxesSubplot: >
```



```
In [110]: data.hist(bins=10,color='green', figsize=(8,6))  
plt.tight_layout()  
plt.show()
```




```
In [111]: plt.scatter(data['TV'], data['Radio'], c=['green' if length>=100 else 'red' for length in data['length']])  
plt.xlabel("TV")  
plt.ylabel("Radio")  
plt.show()
```






```
In [79]: sns.heatmap(corr_matrix, annot = True)  
plt.title("Correlation matrix")  
plt.show()
```



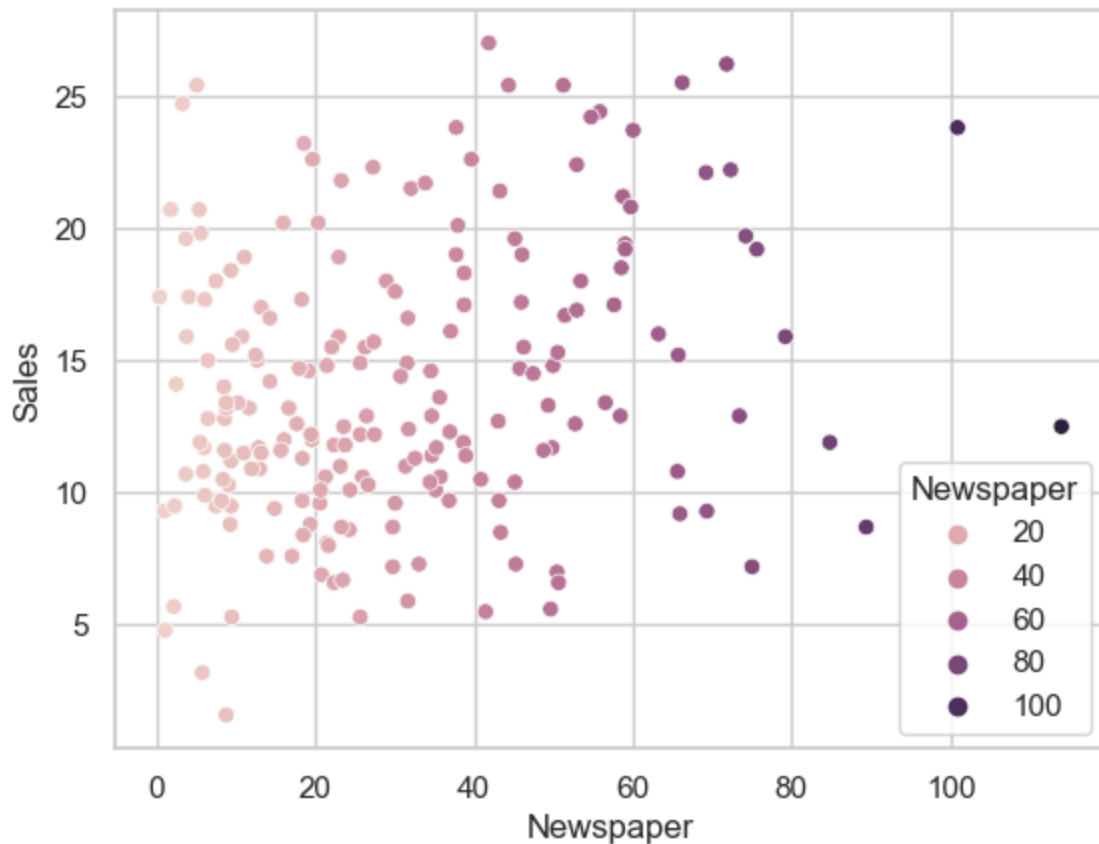
```
In [80]: sns.scatterplot(x='TV', y='Sales', hue='TV', data=data)
```

```
Out[80]: <AxesSubplot: xlabel='TV', ylabel='Sales'>
```



```
In [81]: sns.scatterplot(x='Newspaper', y='Sales', hue='Newspaper', data=data)
```

```
Out[81]: <AxesSubplot: xlabel='Newspaper', ylabel='Sales'>
```



## Training Model

```
In [82]: # Linear Regression
```

```
In [83]: model = LinearRegression()
```

```
In [84]: model.fit(x_train, y_train)
```

```
Out[84]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with [nbviewer.org](https://nbviewer.org).**

```
In [85]: y_predict = model.predict(x_test)
```

```
In [86]: # Evaluation
```

```
In [87]: # Mean Squared Error
rmse = mean_squared_error(y_test, y_predict, squared = False)
rmse
```

Out[87]: 1.78159966153345

```
In [89]: r2 = r2_score(y_test, y_predict)
r2
```

Out[89]: 0.899438024100912

```
In [93]: # XGBRegression

from xgboost import XGBRegressor
regressor = XGBRegressor()
regressor.fit(x_train, y_train)
```

Out[93]: XGBRegressor(base\_score=None, booster=None, callbacks=None, colsample\_bylevel=None, colsample\_bynode=None, colsample\_bytree=None, early\_stopping\_rounds=None, enable\_categorical=False, eval\_metric=None, feature\_types=None, gamma=None, gpu\_id=None, grow\_policy=None, importance\_type=None, interaction\_constraints=None, learning\_rate=None, max\_bin=None, max\_cat\_threshold=None, max\_cat\_to\_onehot=None, max\_delta\_step=None, max\_depth=None, max\_leaves=None, min\_child\_weight=None, missing=nan, monotone\_constraints=None, n\_estimators=100, n\_jobs=None, num\_parallel\_tree=None, predictor=None, random\_state=None, ...)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [100]: # Prediction of training data
tdata_predict = regressor.predict(x_train)
```

```
In [101]: # R-Squared value
r2_test = r2_score(y_train, tdata_predict)
r2_test
```

Out[101]: 0.9999999400896089

```
In [ ]: # Prediction of training data
darta_predict = regressor.predict(x_train)
```

```
In [102]: # R-Squared value
r2_test = r2_score(y_train, rdata_predict)
r2_test
```

Out[102]: 0.9999999400896089

```
In [104]: # Random Forest Regression
random = RandomForestRegressor(random_state=42)
random.fit(x_train, y_train)
```

Out[104]: RandomForestRegressor(random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [105]: # Predicting on test set
y_predict = random.predict(x_test)
```

```
In [106]: # Mean squared_error
rmse = mean_squared_error(y_test, y_predict, squared = False)
rmse
```

Out[106]: 0.7685910811348248

```
In [107]: r2 = r2_score(y_test, y_predict)
r2
```

Out[107]: 0.9812843792541843